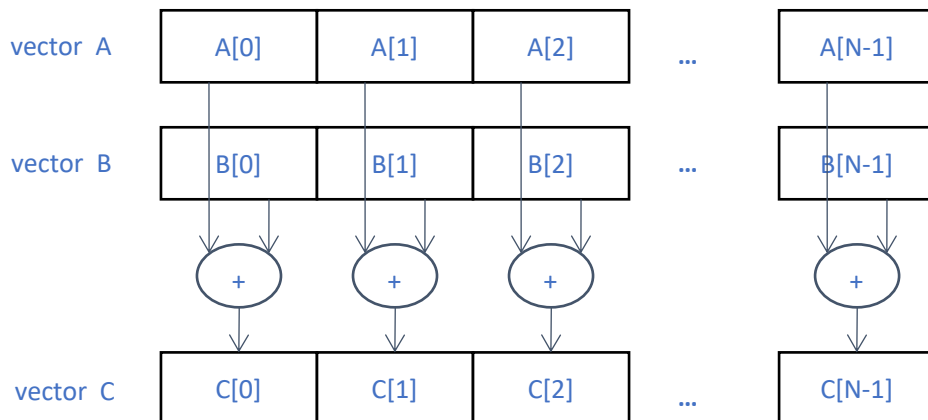


Programming Practice 1: Vector Addition

Objective: To understand basic programming model of GPU-accelerated heterogeneous computing.

Vector addition:

- Given two arrays on n integers A, B.
- An empty array C of size n .
- Add elements of A and B in corresponding positions (i.e., compute $C[i] = A[i] + B[i]$).



(NVIDIA and UIUC, 2017)

Simple flow in GPU-accelerated heterogeneous computing:

- Part 1: Copy input data from CPU memory to GPU memory.
- Part 2: Load GPU program and execute, caching data on GPU chip for full performance.
- Part 3: Copy results from GPU memory to CPU memory.

Host code template:


```
#include<cuda.h>

int main()
{
    // allocate and initialize host memory
    int n = 512; int *h_a = ..., *h_b = ...; *h_c = ...(empty);

    // Part 1
    // allocate device memory for a, b, and c
    // copy a and b to device memory

    // Part 2
    // kernel launch code which let the device performs the actual vector addition

    // Part 3
    // copy c to host memory
    // free device memory
}
```



(NVIDIA and UIUC, 2017)

Basic CUDA Device Memory Management API functions:

- *cudaMalloc()*: allocates an object in the device global memory with the following two parameters in order:
 - Address of a pointer to the allocated object.
 - Size of allocated object in terms of bytes.

Example: `int size = n * sizeof(float); float *d_A;
cudaMalloc((void **) &d_A, size);`

- *cudaFree()*: frees object from device global memory with the following parameter.
 - pointer to freed object.

Example: `cudaFree(d_A);`

- *cudaMemcpy()*: memory data transfer with the following four parameters in order:
 - pointer to destination
 - pointer to source
 - number of bytes copied
 - Type/direction of transfer

Example: `cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);`

Practice 1.1: Implement host code for parallel addition of 512-element vectors.

Practice 1.2: Implement device code for parallel addition of 512-element vectors.

Practice 1.3: Combine your codes in 1.1 and 1.2 to compute parallel addition of 512-element vectors.

Practice 1.4: Adjust your codes in 1.3 to compute parallel addition of n -element vectors.